

Decision trees

Terence Parr
MSDS program
University of San Francisco

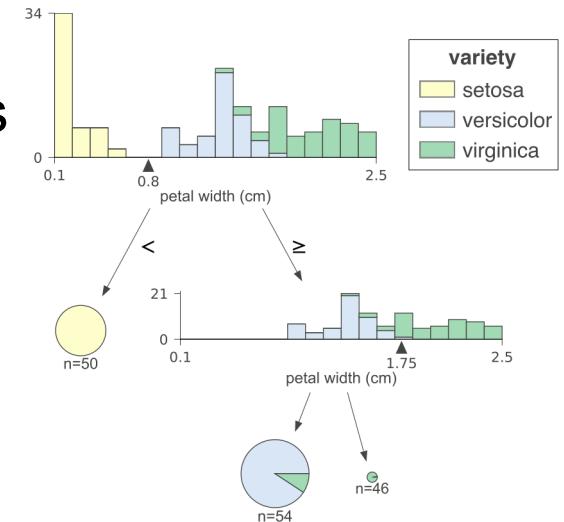


The essence of decision trees

- Decision trees are like kNN , but with rectangular not polygonal hypervolumes and dynamic k not fixed k
- Partition feature space into tight rectangular hypervolumes with y values as pure/similar as possible for records in that hypervolume
- Not so tight that the hypervolumes have too few vectors (records/samples), which tends to overfit the training data
- Prediction for unknown vector:
 - predict the mean y for training samples in that region (regression)
 - predict the mode (most common) y in that region (classification)
- Binary trees just happen to be an efficient implementation

Basic properties of decision tree models

- Decision trees consist of internal decision nodes and leaf nodes that make predictions
- Each input record (feature vector) associated w/exactly one leaf
- Each leaf has 1 or more records whose y's are as pure as possible (model hyperparameters change affect number of records per leaf)
- Prediction proceeds from root to leaf, testing var/value combos
- Regressor: leaf predicts average of y for associated records
- Classifier: leaf predicts mode (most common) class

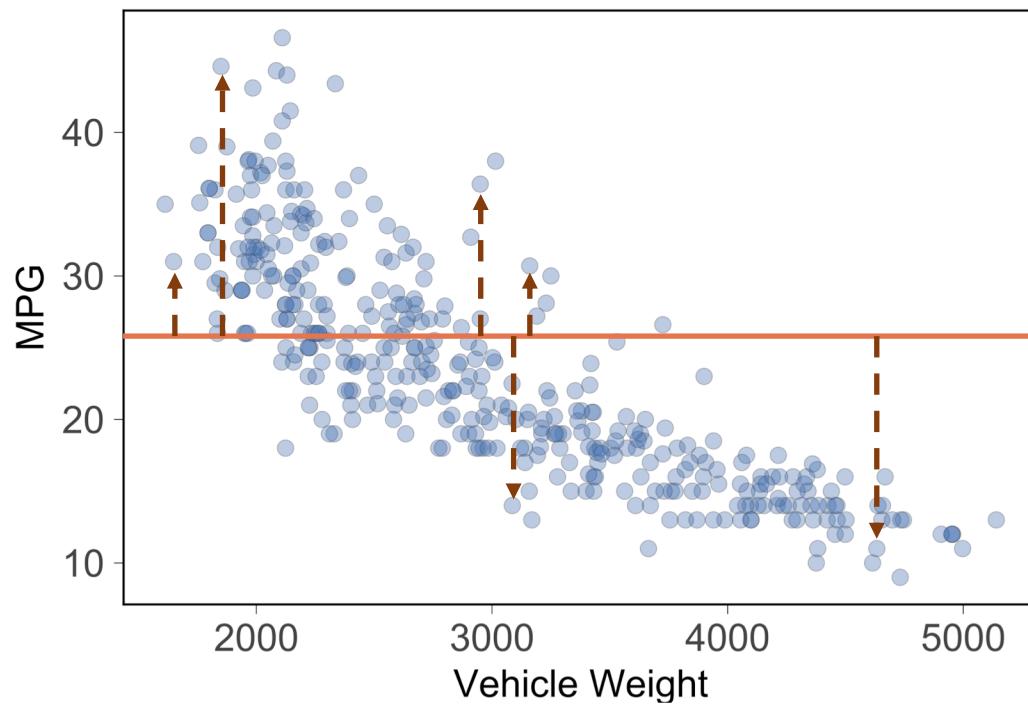


Most images in these slides from package <https://github.com/parrt/dtreeviz>

Let's reinvent decision trees

Let's create a simple regressor in 1D

Predict MPG
from weight



Predict a single
constant for
entire region
(the mean)

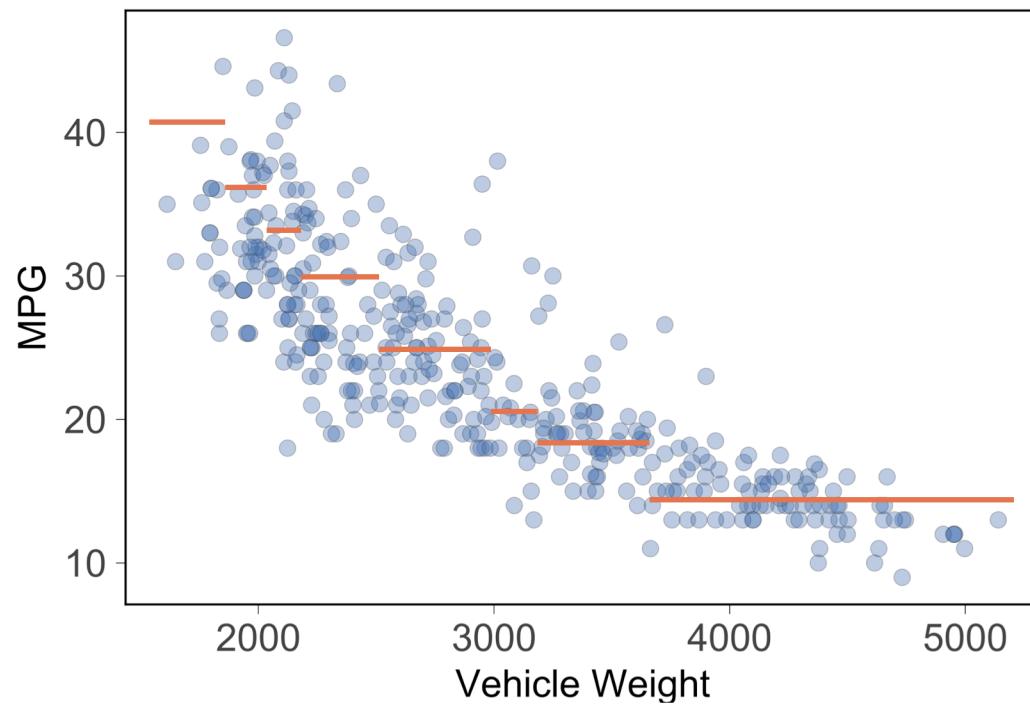
MSE is high

↑ Residuals from
| point to region
mean

Improve by partitioning, using multiple lines

*WHERE DO
WE SPLIT??*

Find groups
of similar MPG
values, which
yields a lower
MSE

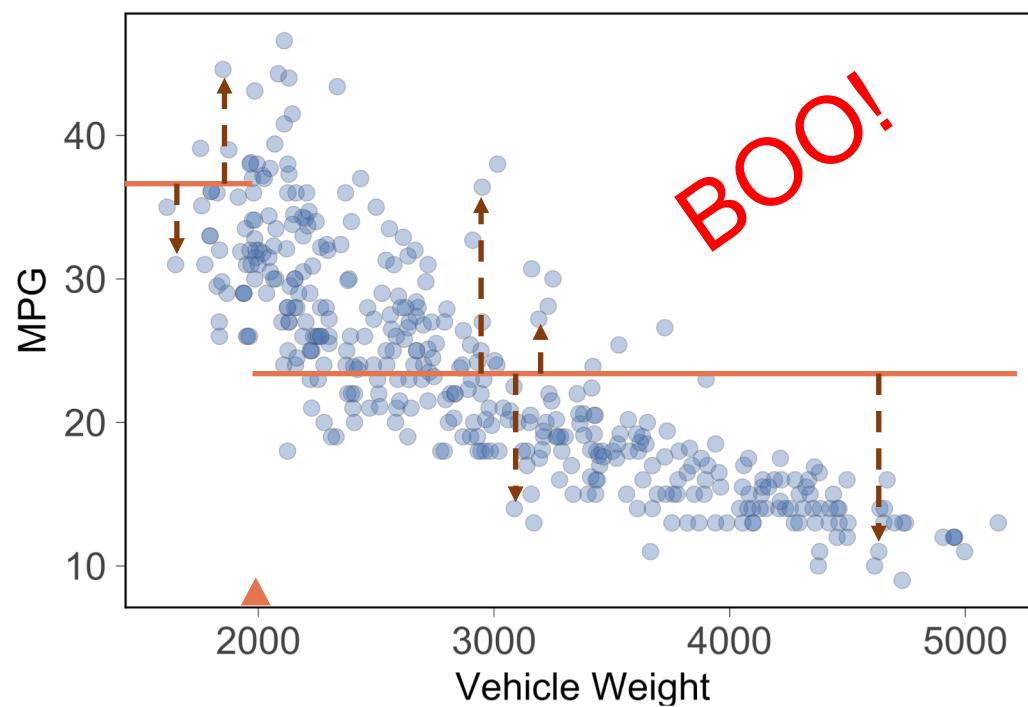


Can only use
horizontal lines,
but can use lots

Each region
predicts mean
in piecewise
fashion

Strategy: find split point giving least MSE

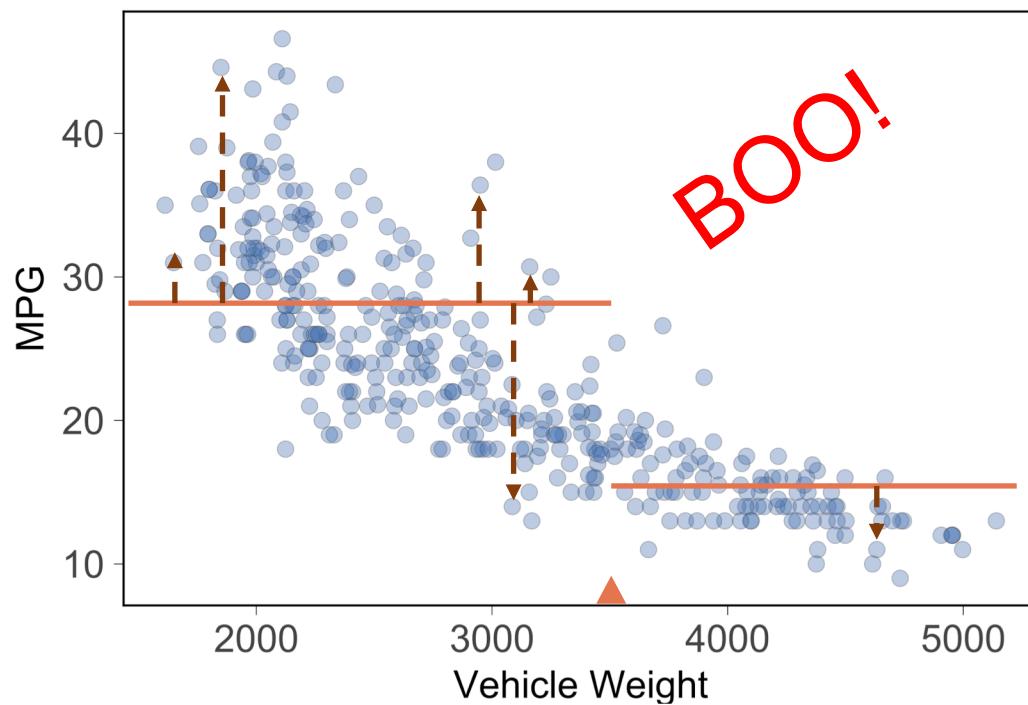
Split region
into two
subregions,
each
predicting
mean



X=2000 is
BAD CHOICE:
MSE very high

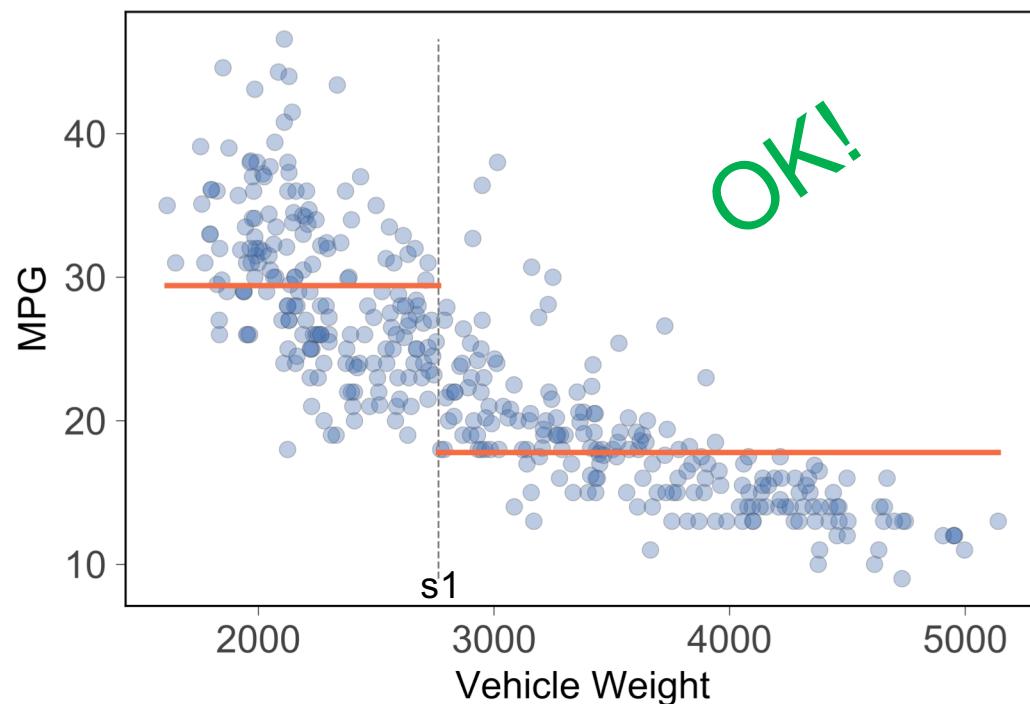
Strategy: find split point giving least MSE

Split region
into two
subregions,
each
predicting
mean



X=3500 is
ANOTHER
BAD CHOICE:
MSE very high

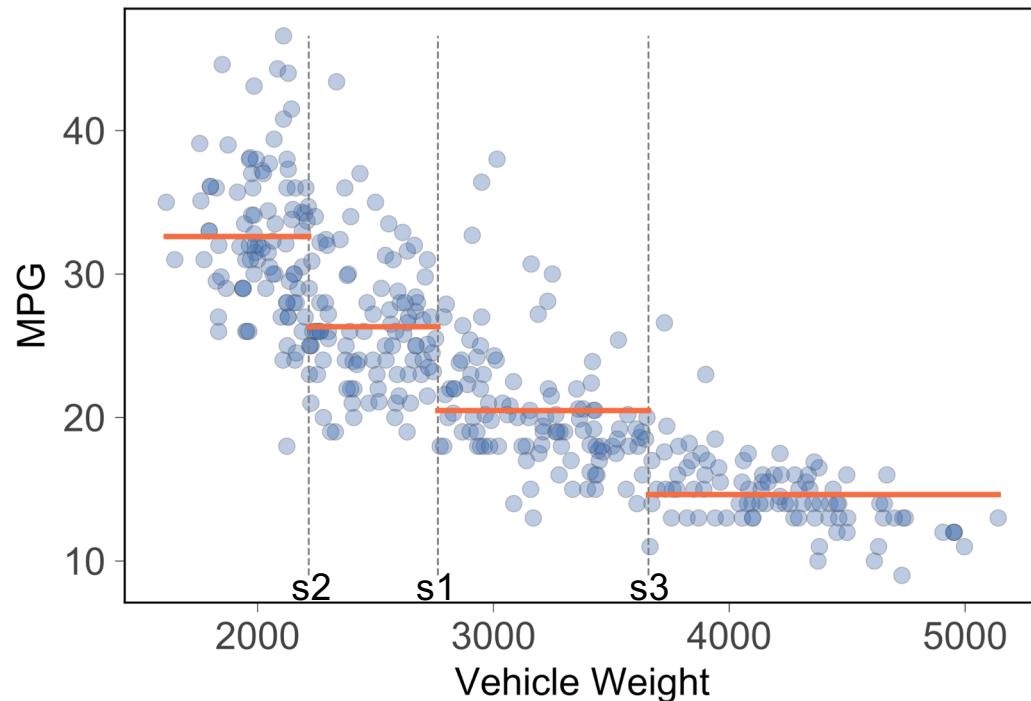
A split exists that gives min MSE for regions



Technique:
Exhaustively
check all feature
values, computing
MSE or variance
for each split

Choose split point
giving min MSE

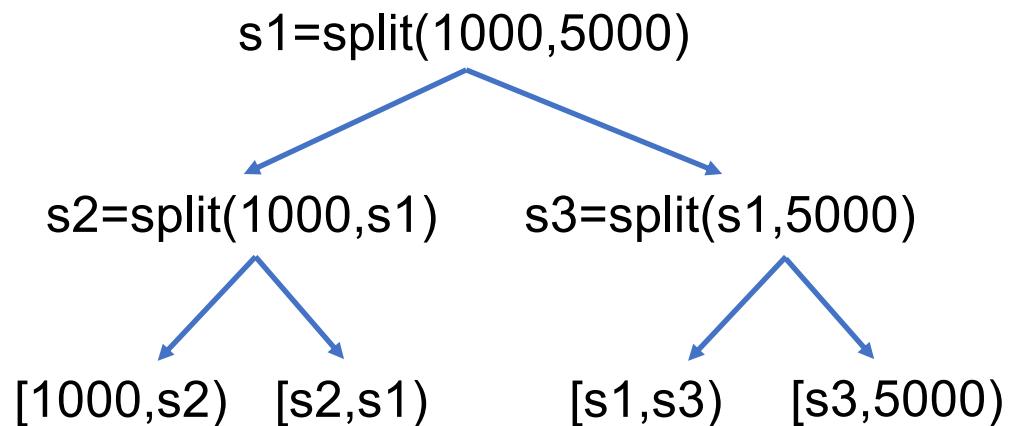
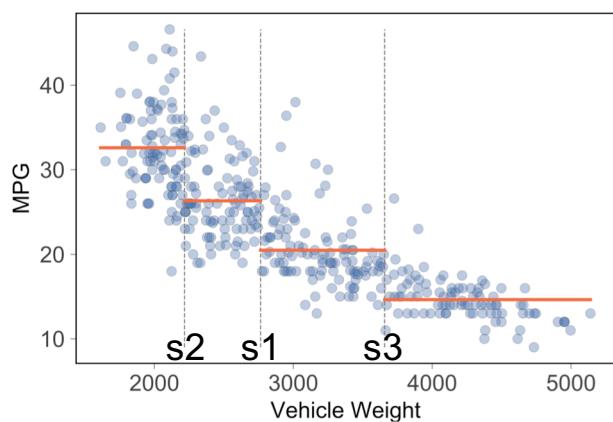
Now split those 2 regions to get 4 regions



Split s_1 stays,
recursively split
left/right regions to
get splits s_2, s_3

Kinda like binary
search or other
divide-and-conquer
strategy

Recursive-call tree from model training gives regions defined by splits s_1, s_2, s_3

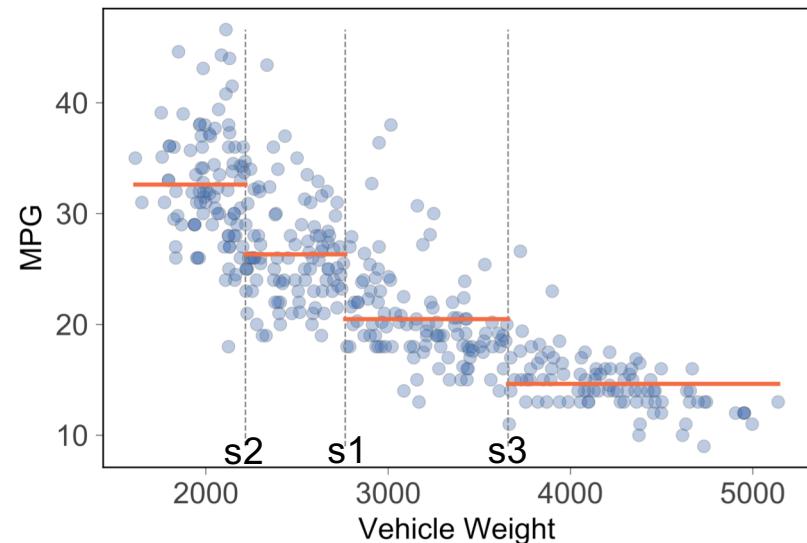


Split (recurse) until:

- All potential splits do not reduce MSE
- All nodes have min num samples
- Max number of splits reached
- Etc...

*Predictions are avg of MPG
(target) values in regions*

Hardcoded non-tree model implementation

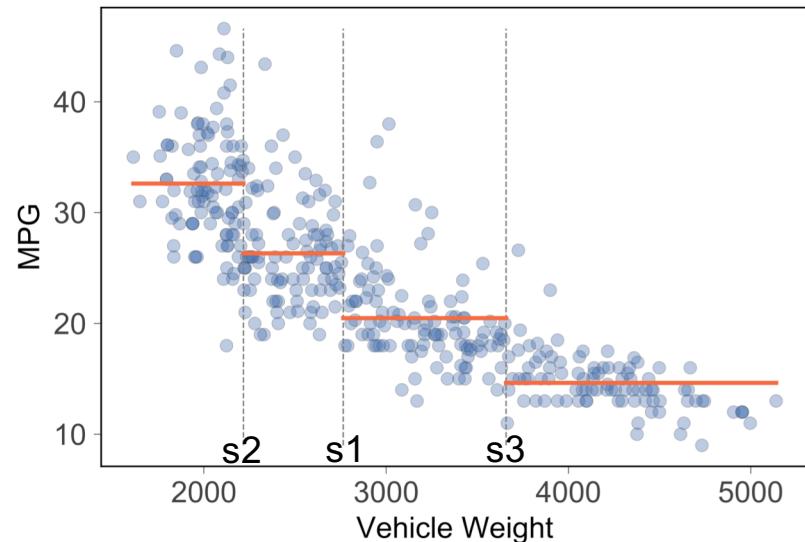


To partition space, test in recursion/split order

if $x < s_1$ and $x < s_2$: predict 32.6
if $x < s_1$ and $x \geq s_2$: predict 26.3
if $x \geq s_1$ and $x < s_3$: predict 20.5
if $x \geq s_1$ and $x \geq s_3$: predict 14.6

Note repeated comparisons!

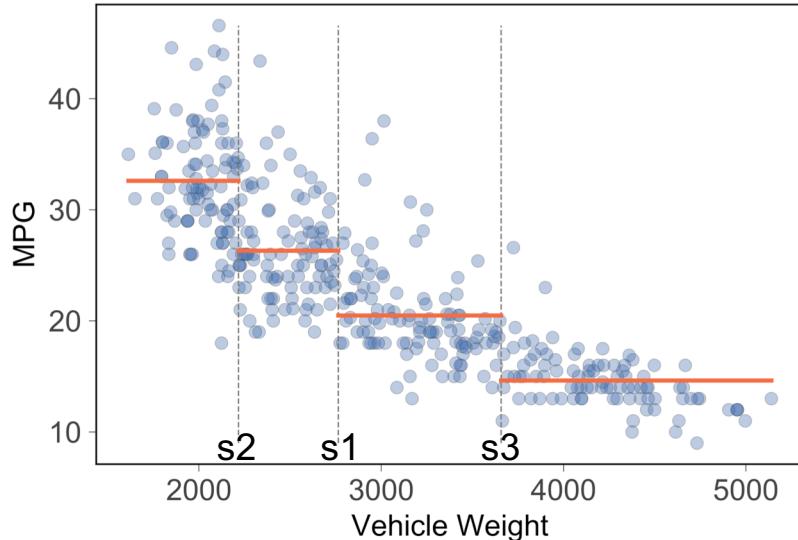
Factor the split comparisons for efficiency



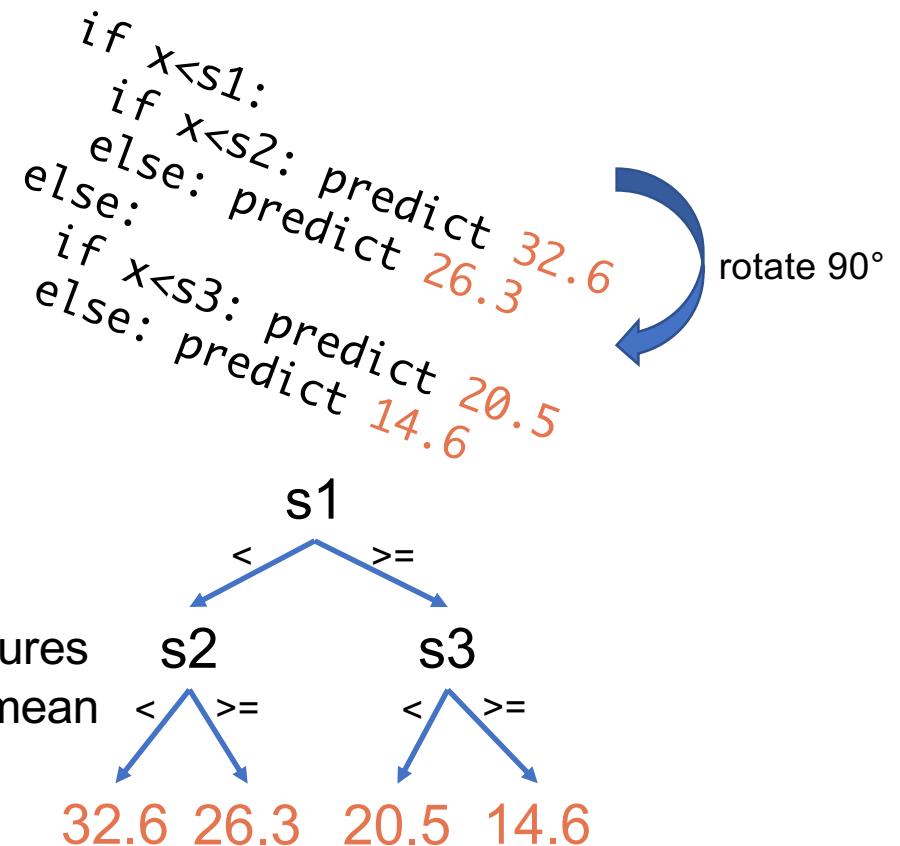
```
if x<s1:  
    if x<s2: predict 32.6  
    else: predict 26.3  
else:  
    if x<s3: predict 20.5  
    else: predict 14.6
```

But, don't want to hardcode model!

Represent nested conditionals as tree



Internal nodes test features
Leaves predict region mean

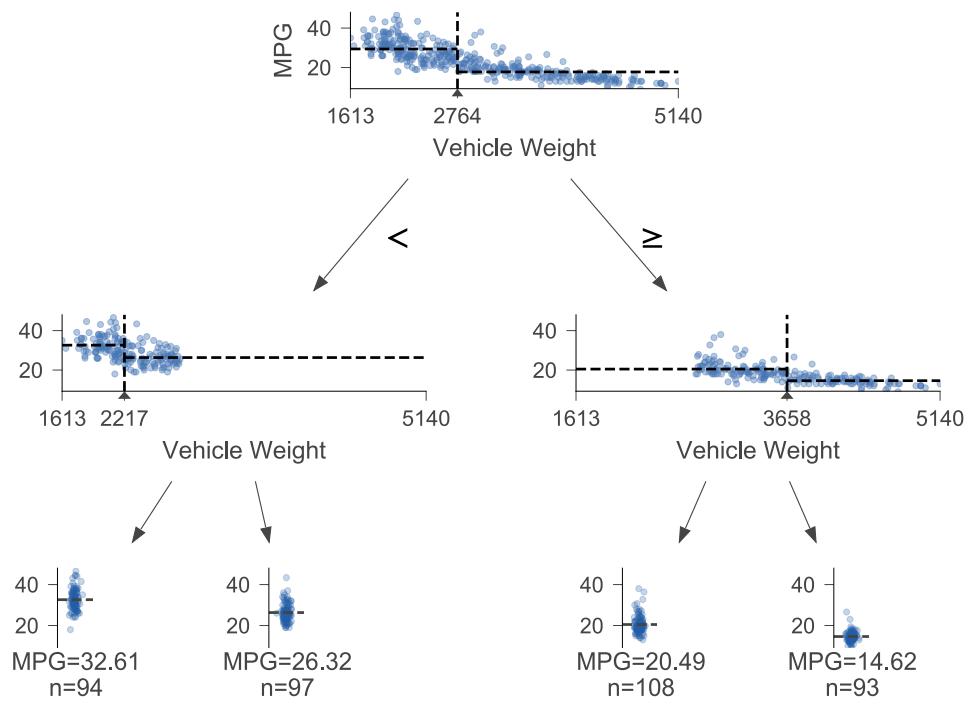
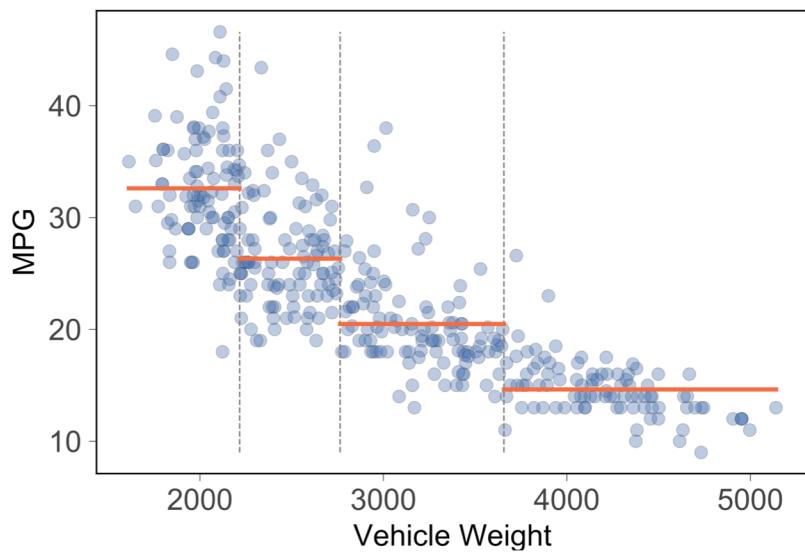


We morph tree of recursion from training into decision tree!



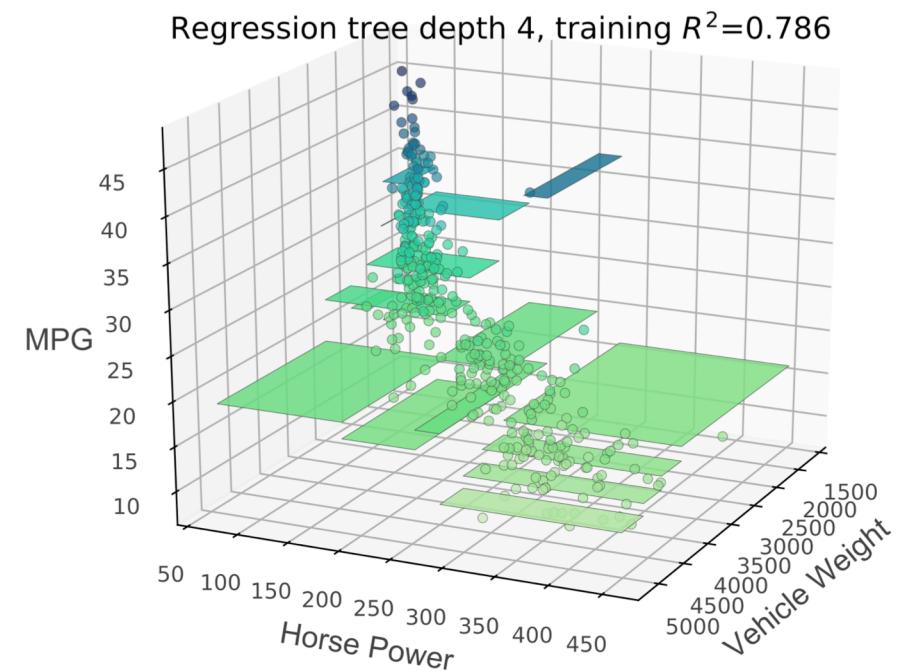
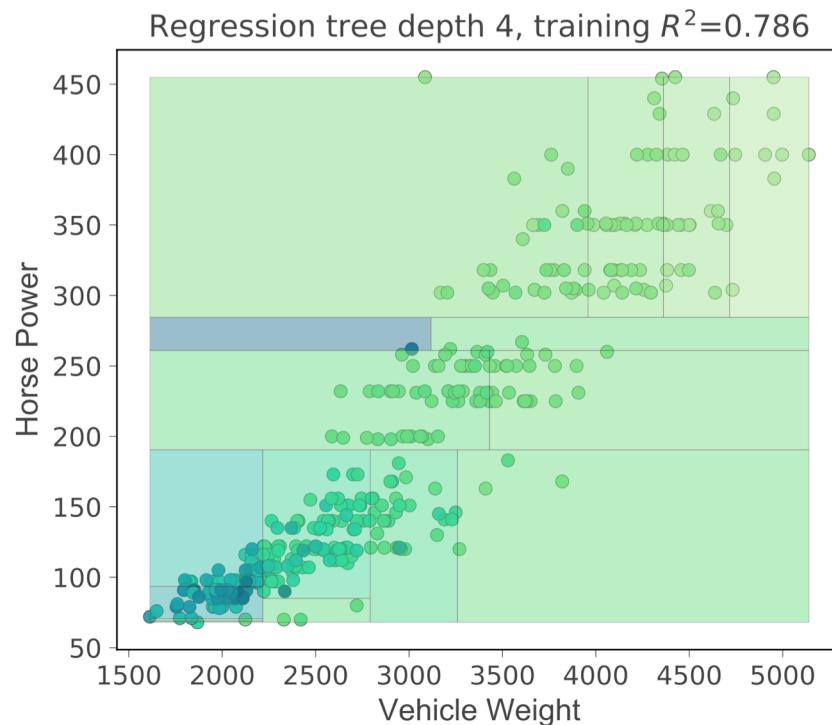
UNIVERSITY OF SAN FRANCISCO

1D feature space vs dtreeviz decision tree

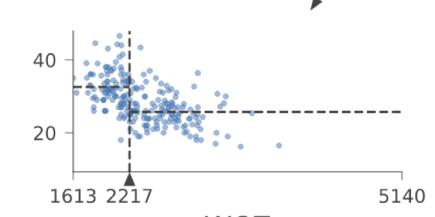
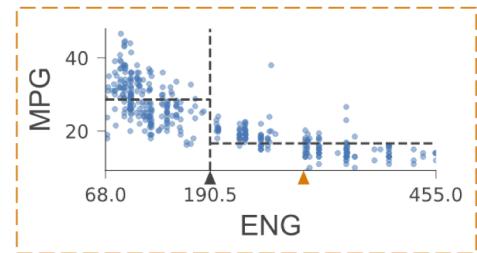


<https://github.com/parrt/msds621/blob/master/notebooks/trees/partitioning.ipynb>

2D regressor feature space (heatmap, 3D)

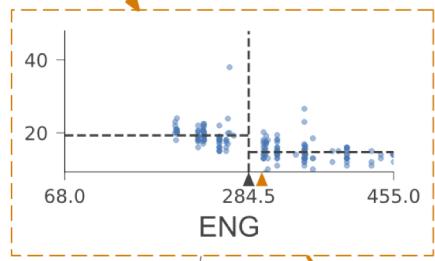


How predictions are made from unknown x



MPG=32.61
n=94

MPG=25.73
n=128

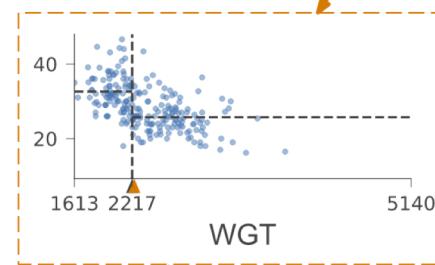
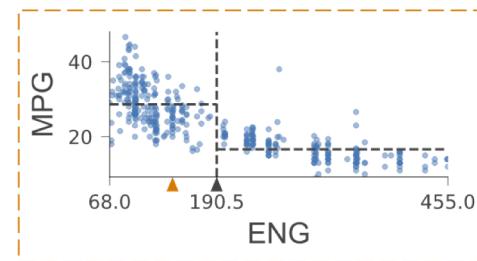


MPG=19.32
n=72

MPG=14.71
n=98

Prediction 14.71

ENG	WGT
300.00	3500.00

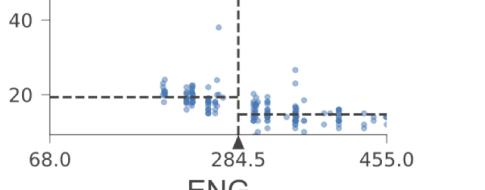


MPG=32.61
n=94

MPG=25.73
n=128

Prediction 25.73

ENG	WGT
140.00	2240.00



MPG=19.32
n=72

MPG=14.71
n=98

An aside: partitioning, tree viz done with custom library

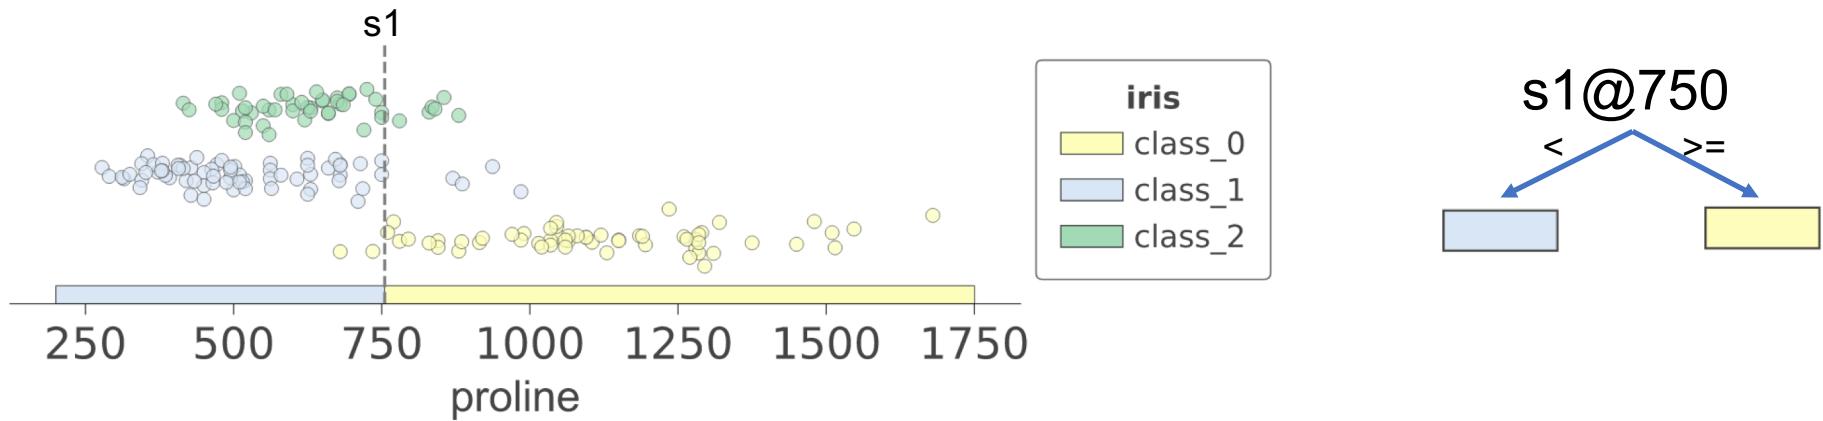
- Do “`pip install dtreeviz`”
- I built with Prince Grover, previous MSDS student
- See <https://github.com/parrt/dtreeviz> and the article for more detail: <https://explained.ai/decision-tree-viz/index.html>
- Advice: never accept status quo; always strive for more / better
- See “*How to lead a fulfilling life by being dissatisfied*” buried in my talk on decision tree viz
https://twitter.com/the_antlr_guy/status/1120359898062000128

Classifiers

Classifiers split feature space too

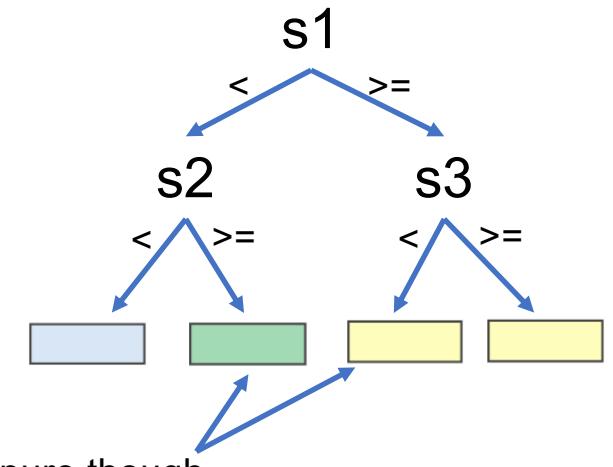
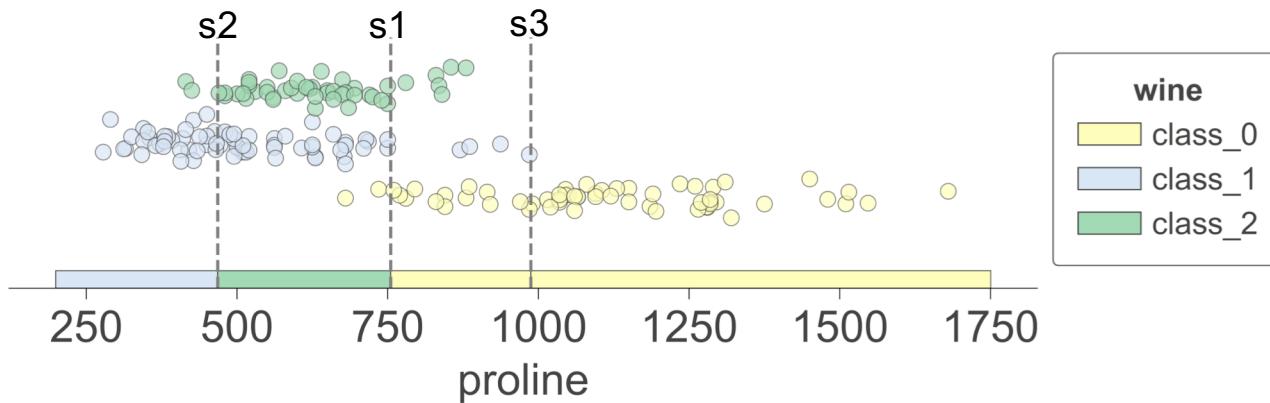
Predict wine
from proline

- Internal decision nodes test features just like regressor trees
- Leaves predict most common target category (mode) not mean
- Find split that decreases average impurity of left/right subregions

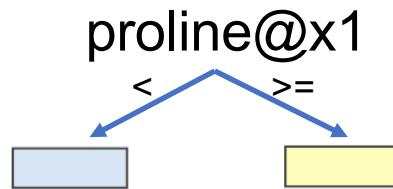


<https://github.com/parrt/msds621/blob/master/notebooks/trees/partitioning.ipynb>

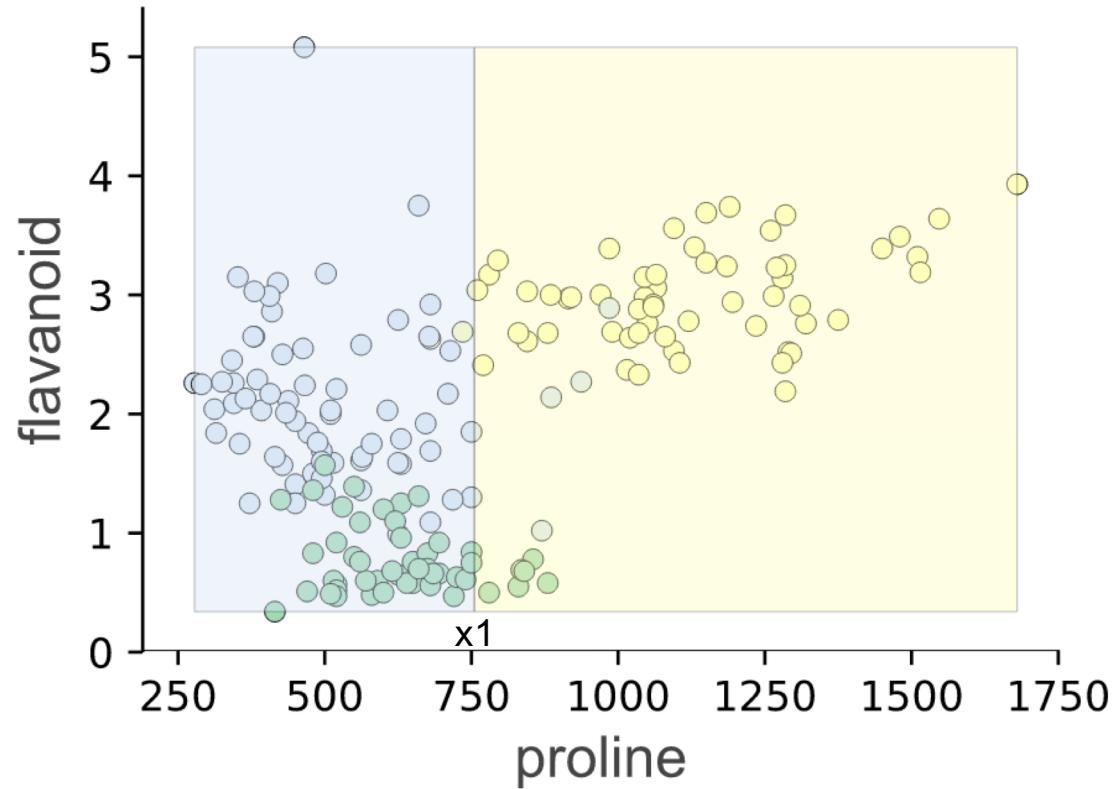
Split s1 subregions into more subregions



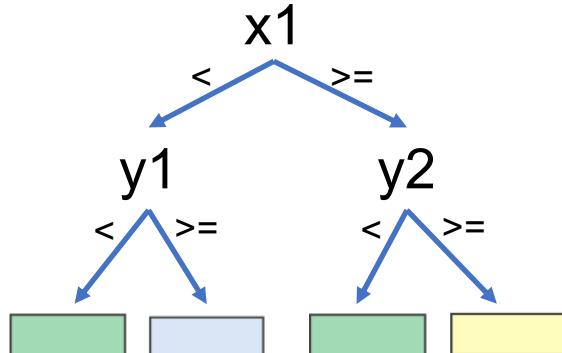
To improve predictions:
Use 2 features and
split 2D feature
space into regions



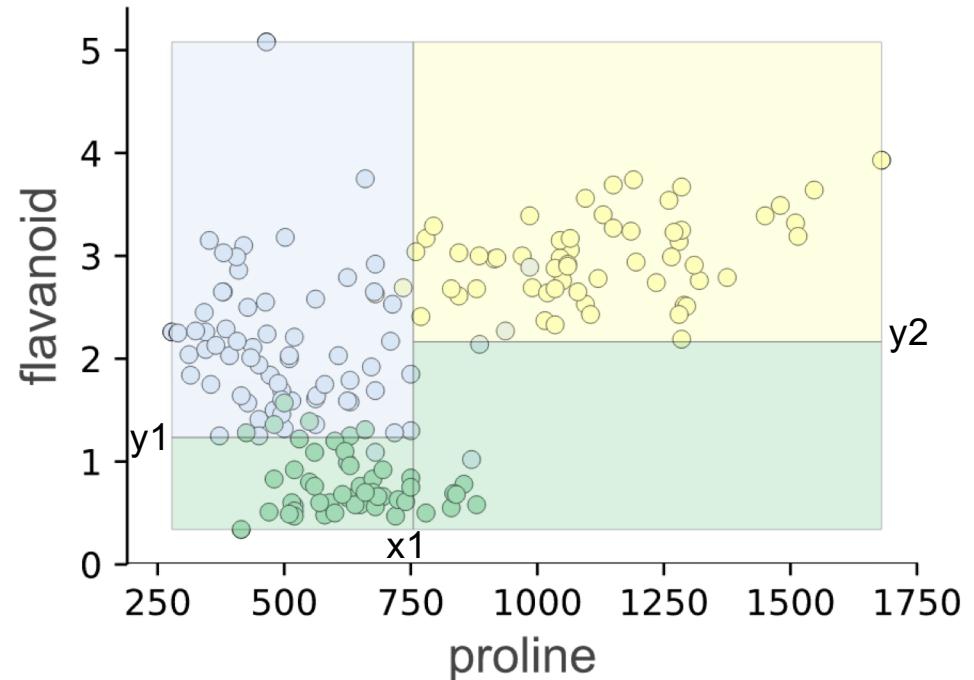
*Training looks for (feature, split point) combos giving more pure subregions.
Decision nodes compare single feature value in vector to split point*



Improve predictions:
Use 2 features and
split 2D feature
space into regions



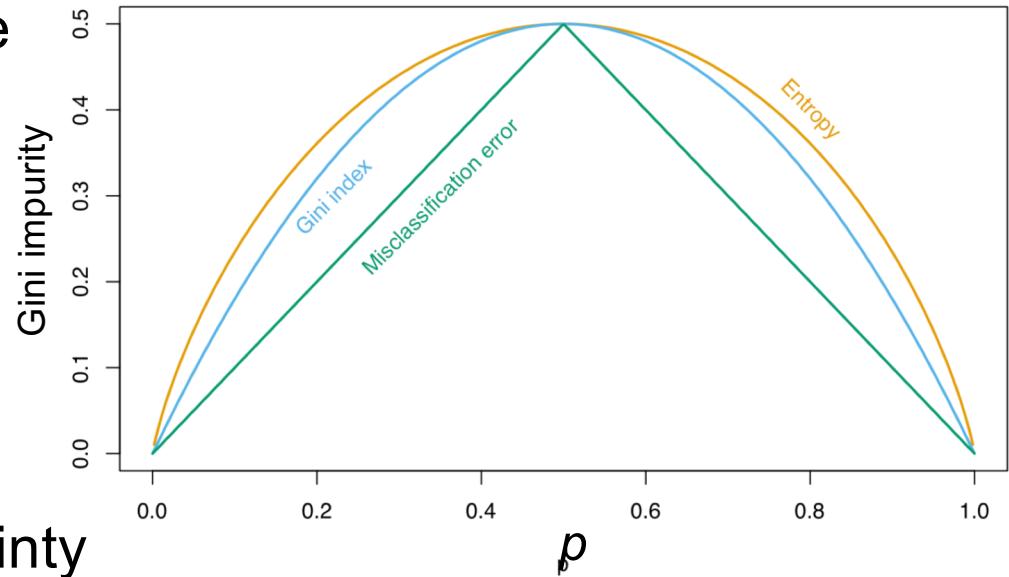
*Training looks for (feature, split point)
 combos giving more pure subregions.
 Decision nodes compare single
 feature value to split point*



How do we measure impurity?

Gini impurity

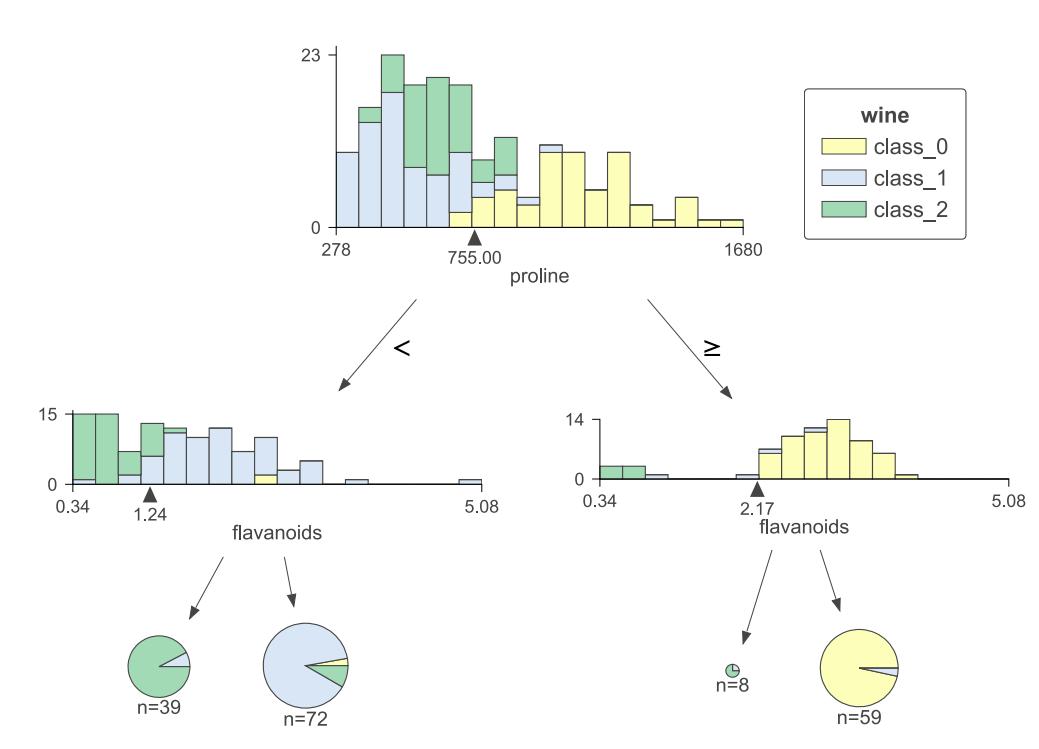
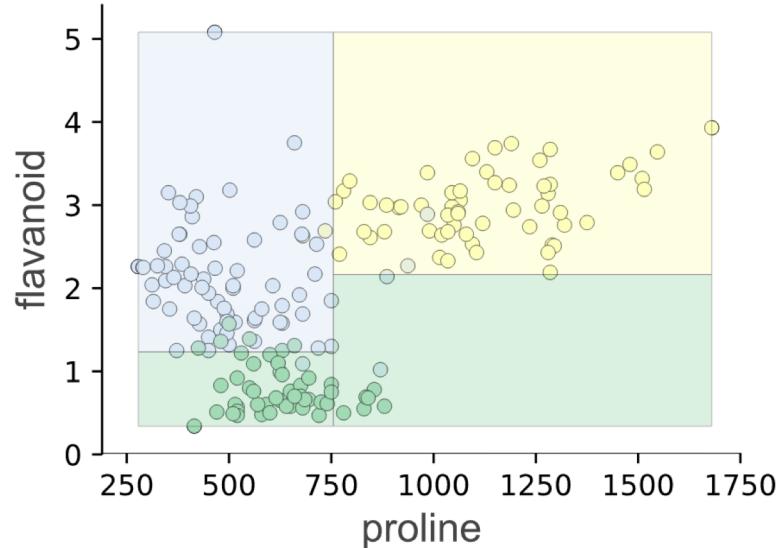
- A measure of y 's uncertainty, like entropy
- Minimize during training
- Let p_i be the fraction of y values with class i (likelihood of i), $k = \text{number of unique classes in } y$
- If all p_i same, max (gini) uncertainty
- If all y same, $p_i=0$ or $p_i=1$, gini=0



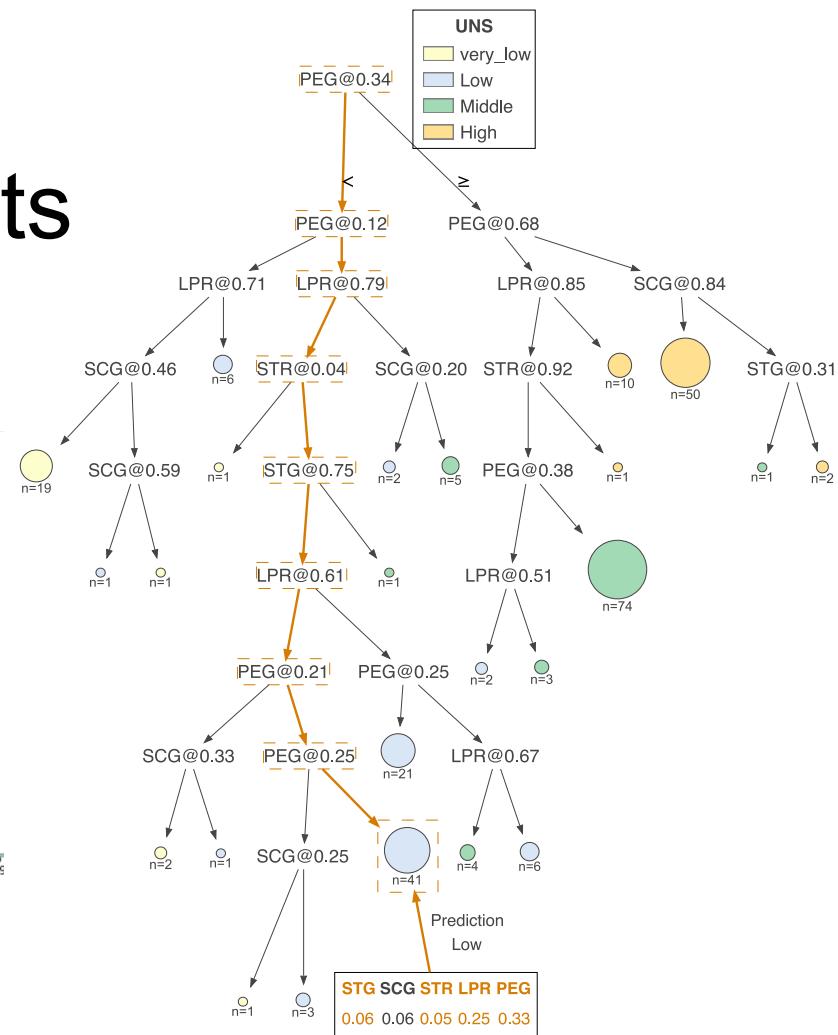
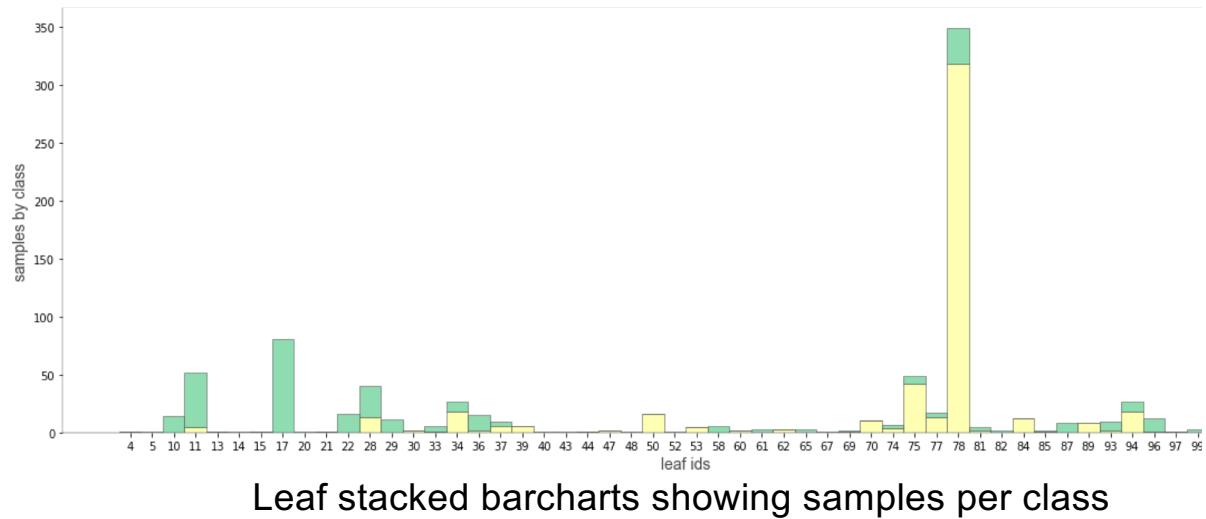
$$Gini(y) = \sum_{i=1}^k p_i \sum_{j \neq i}^k p_j = \sum_{i=1}^k p_i(1 - p_i) = 1 - \sum_{i=1}^k p_i^2$$

https://en.wikipedia.org/wiki/Decision_tree_learning#Gini_impurity

dtreeviz Tree structure for 2D partition



For bigger trees, can
do nonfancy plot, leaf plots



Tree structure's effect on prediction error

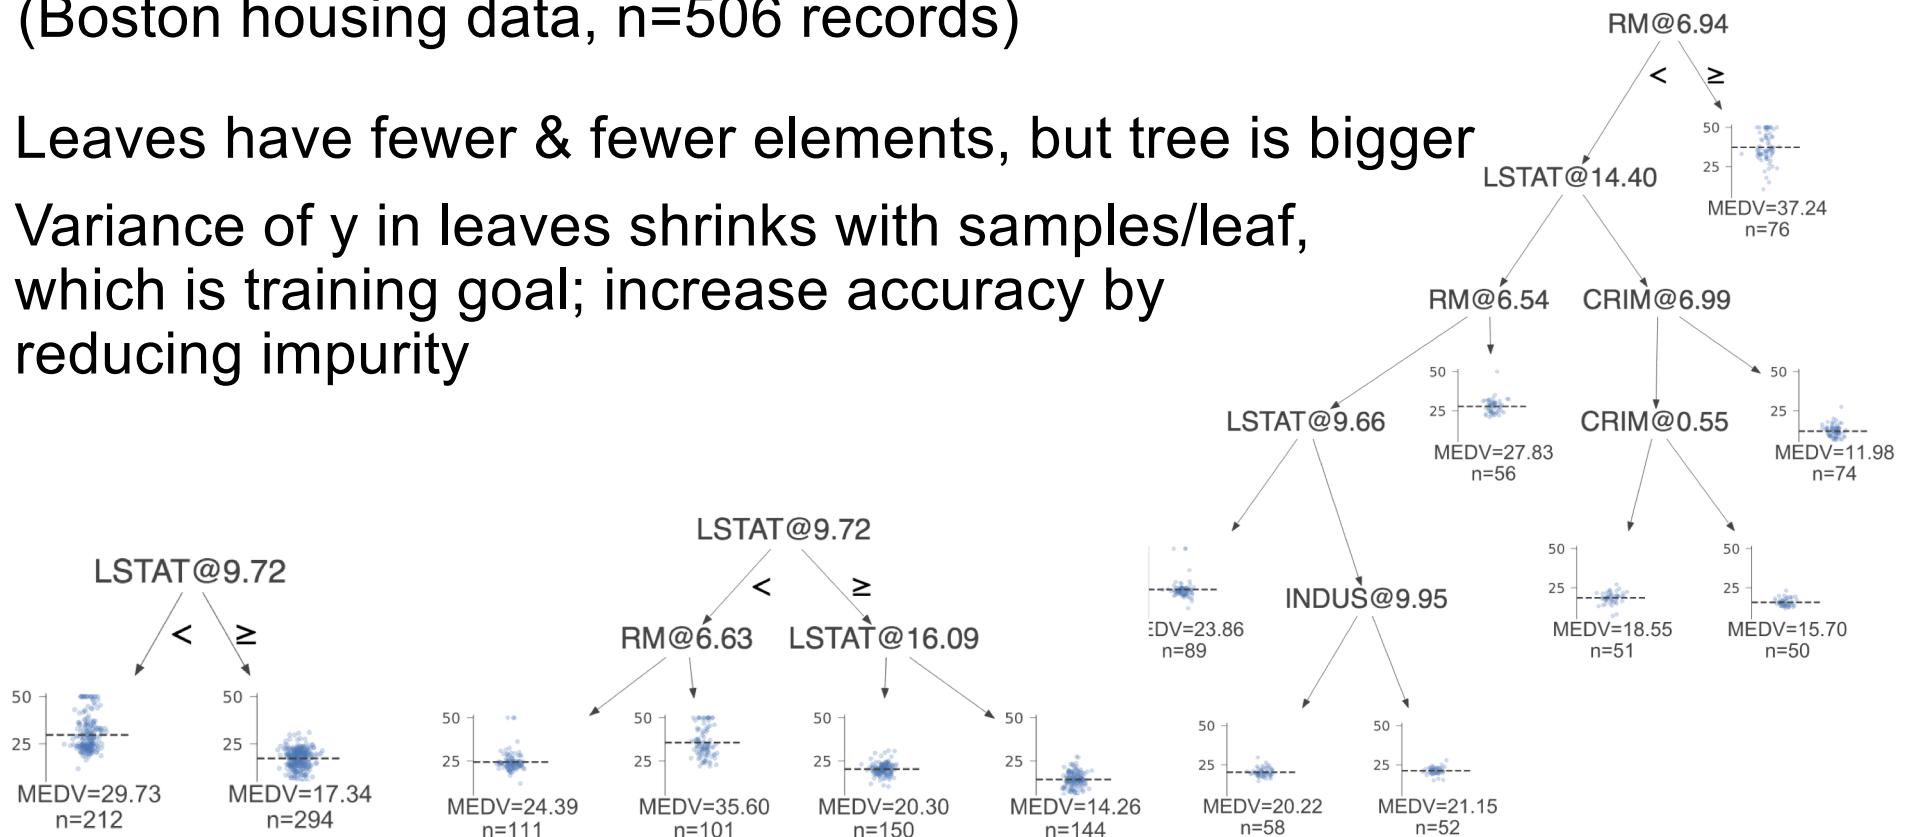
Hyperparameter num_samples_leaf

- Idea: don't split regions w/less than num_samples_leaf records
- Similar to limiting height of tree but finer granularity of control
- More direct control of generality than tree height
- Degenerate case where num_samples_leaf=n (sample size)
 - What does such a regressor predict?
 - What does such a classifier predict?
 - Describe accuracy of this extreme model
 - If we trained on many different training sets pulled from same data distribution, how stable would the test set prediction error be? (What does that say about variance/generality?)

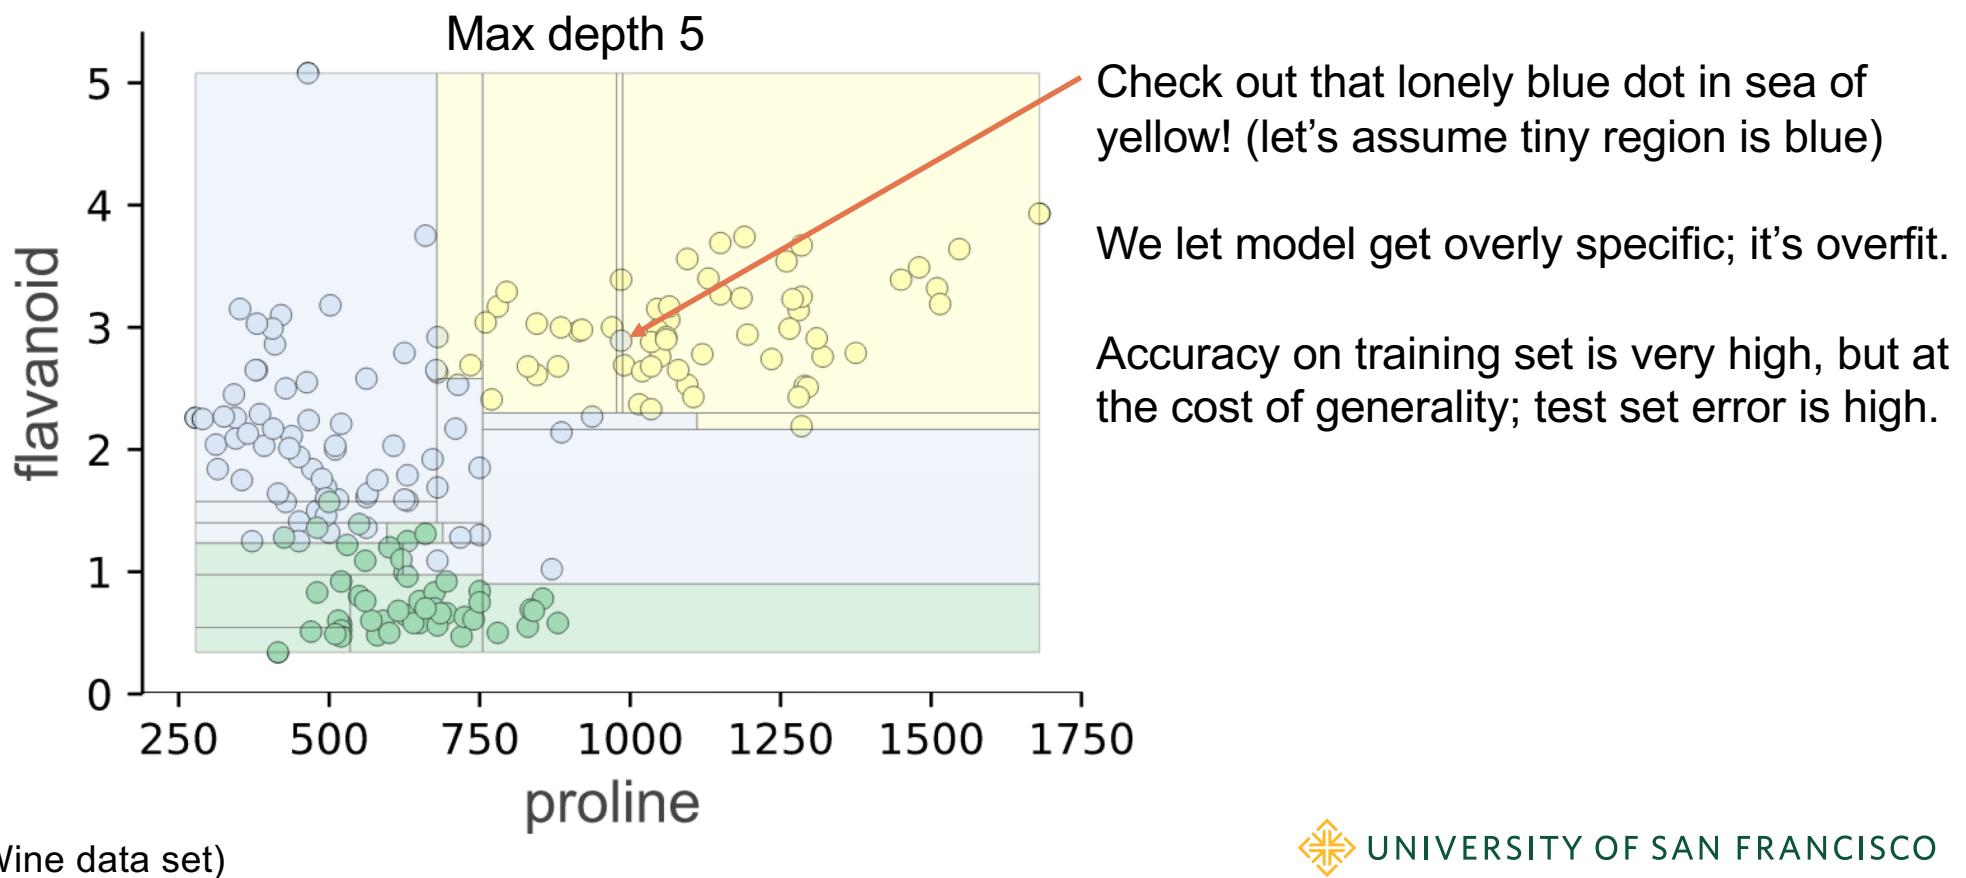
Trees for 200,100,50 samples per leaf

(Boston housing data, n=506 records)

- Leaves have fewer & fewer elements, but tree is bigger
- Variance of y in leaves shrinks with samples/leaf, which is training goal; increase accuracy by reducing impurity

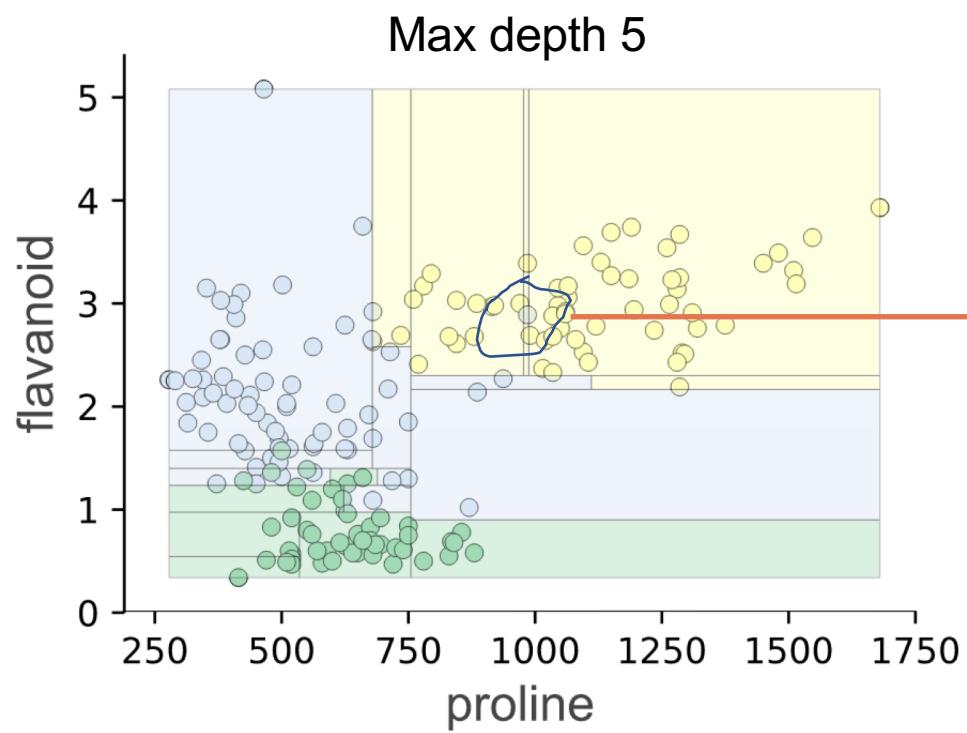


What happens with very small leaves?

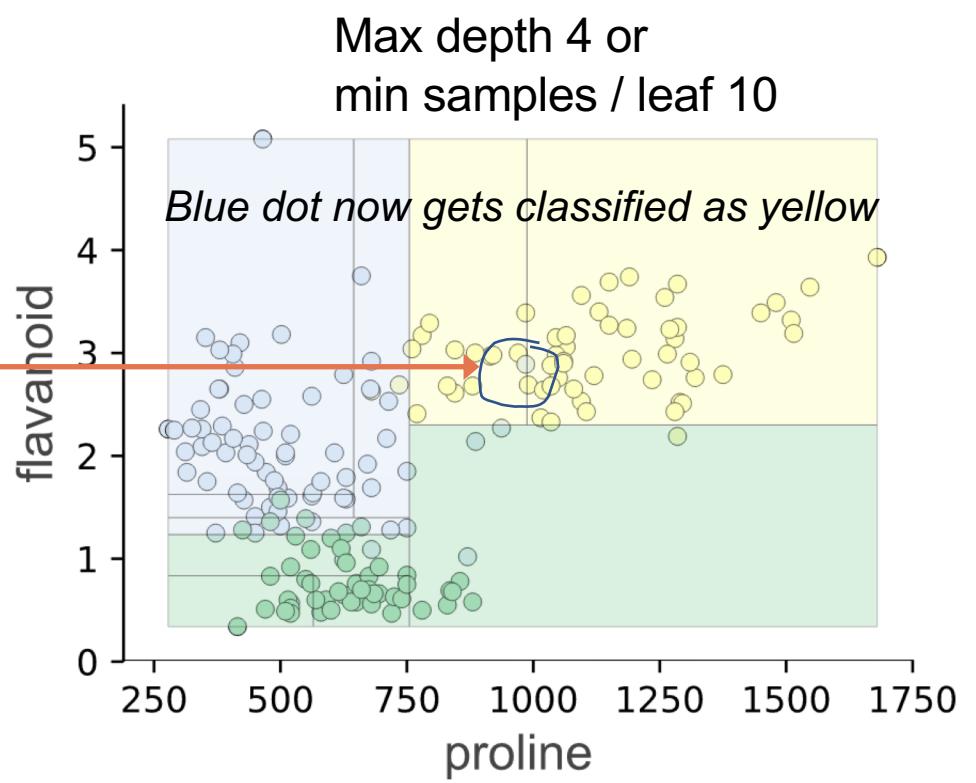


UNIVERSITY OF SAN FRANCISCO

How could we (likely) improve generality?



(Wine data set)



Key takeaways

- Trees partition feature space into tight rectangular hypervolumes with pure/similar y values for records in that hypervolume
- Decision trees have internal decision nodes that test variables at split points and leaf nodes that make predictions
- Leaves predict mean (regressor) or mode (classifier) of samples
- Partitioning subject to reducing MSE or Gini
- Limiting tree height or leaf size increases bias but improves generality (reduces variance)
- (We'll have whole lecture on training these beasts)

Lab time

- Partitioning feature space
<https://nbviewer.jupyter.org/github/parrt/msds621/blob/master/labs/trees/partitioning-feature-space.ipynb>
- Note: that links is for viewing since github screws that up; you will need to go to the course repo to get interactive labs
<https://github.com/parrt/msds621/tree/master/labs/trees>